

NeoCore FAQ

Technical reference for platform integrators, security reviewers, and procurement teams

This FAQ covers the questions enterprise customers most often ask before deploying NeoCore at scale. It is organised across six areas: data security and privacy, transcript quality, API access, reliability and support, device filtering and management, and hardware. Items marked "in progress" are on the active roadmap; specific commercial commitments (SLAs, lead times, pricing tiers) are agreed under the master service agreement.

1. Data Security and Privacy

Is NeoCore compliant with the Digital Personal Data Protection Act, 2023 (DPDPA)?

NeoSapient is ISO 27001 certified for information security management, and every third-party service in the platform stack meets SOC 2 compliance standards. Customer data is held on India-resident servers.

In progress — Full DPDPA compliance work is in active progress, with formal compliance expected within 6 to 8 weeks. For deployments processing personal data, NeoSapient acts as a Data Processor on behalf of the customer (Data Fiduciary); a Data Processing Agreement is executed alongside the master service agreement before production data flows through the platform.

Where is conversation and transcript data stored?

Production data is stored on India-resident servers. This makes NeoCore suitable for deployments where data residency in India is a regulatory or contractual requirement, including healthcare, financial services, and government use cases.

Is raw audio retained on NeoSapient servers?

No. Raw audio is not retained by default — it is processed transiently for transcription and discarded once the transcript and structured outputs are generated. Where a customer has a specific audit or quality-review need for short-term audio retention, this can be configured as a separately agreed arrangement.

What is the data retention policy for transcripts?

Transcripts are retained indefinitely unless a deletion request is raised. Customers requiring a defined retention window (e.g., 90 days, 1 year, 7 years) for their own regulatory alignment can request this as part of the master service agreement, and it is implemented via lifecycle policies on the tenant.

Is customer data used to train NeoSapient's models?

No. Customer conversation data is not used for model training or product improvement.

How are personally identifiable information and personal conversations handled?

PII and personal (non-work) conversations are detected and redacted as part of the processing pipeline. This reduces the volume of sensitive content stored in transcripts and limits the scope of personal data flowing through downstream customer systems.

How is data encrypted and access controlled?

All API traffic is over TLS. The NeoCore API uses JWT-based authentication with token rotation and scoped access for every API consumer, so customers can issue scoped credentials per integration, per environment, and per service principal. Tokens can be rotated by the customer at any time.

How often is the platform audited?

NeoSapien commissions independent vulnerability testing and penetration audits on a fixed bi-annual cadence (every six months). Audit summaries and the SOC 2 status of platform sub-processors are available to enterprise customers under NDA.

Can a specific memory or all data for a user be deleted on request?

Yes. A memory deletion endpoint supports granular deletion to honor data principal rights requests. Bulk deletion at the user or tenant level is supported via the administrative API for enterprise tenants. On termination of the contract, customer data is deleted in line with the timelines specified in the DPA.

How are data breaches handled?

In the event of a confirmed personal data breach affecting customer data, NeoSapien notifies the customer without undue delay. The customer remains responsible for any onward notification to regulators and data principals under applicable law; NeoSapien provides the information and assistance required to support that process. Specific notification timelines are set out in the DPA.

2. Transcript Quality and Content

What is the latency from speech to a synthesised, retrievable memory?

Current end-to-end latency for memory synthesis runs at approximately 2 minutes at the P50 (median) and approximately 10 minutes at the P99. This window covers transcription, diarization, and synthesis. Customers building latency-sensitive workflows (live alerting, in-conversation prompts) should size their integration around the P99 rather than the median. Sub-minute latency is on the development roadmap; deployments today should not assume it.

Does the memory object include the full raw transcript?

Yes. Each memory object includes the full raw transcript.

How accurate is speaker diarization in multi-speaker conversations?

Diarization accuracy is 87%+ for multi-agent conversations with a clear audio stream. In environments with significant ambient noise (moving vehicles, busy retail floors, public spaces), accuracy reduces by approximately 4 percentage points. Deployments in noisy environments should be validated during a pilot.

What languages does the platform support?

NeoCore supports 100+ languages with significant coverage of Indic languages. Language-based code-switching is supported — speakers can mix languages within a conversation (e.g., English and Hindi) and the transcription

handles the transitions. This matters in Indian enterprise field environments where multilingual conversations are common.

How does the platform handle domain-specific vocabulary (drug names, SKUs, brand names)?

NeoCore has not been benchmarked on domain-specific lexicons (e.g., medical drug names) and does not promise incremental accuracy over the standard transcription baseline. Where a customer has a domain catalogue — a drug list, product SKUs, brand and competitor names — this can be supplied to NeoSapien and the relevant keywords are boosted in the model. The exact uplift varies by domain and should be validated during a pilot.

Is there a minimum audio duration for a memory to be created?

Yes. Very short utterances are not captured as memories — the platform applies a minimum-duration filter to suppress incidental sounds and false triggers. Customers expecting to capture short verbal callouts (single-word commands, brief acknowledgements) should validate that the relevant audio segments meet the minimum threshold.

Are the AI Outputs suitable for clinical or regulated decision-making?

No. Transcripts, summaries, and synthesised outputs are provided on an "as is" basis and are not medical-grade. They should not be relied upon as a substitute for human review in clinical, diagnostic, or therapeutic decisions, or for other regulated workflows. Customers are responsible for reviewing and validating outputs before use in any decision with regulatory or safety implications.

3. API Access

Is there a developer or enterprise API tier?

Yes. The NeoCore developer API is live and documented. It plugs into existing CRM, BI, and analytics stacks. Enterprise engagements include a stable API contract, sized rate limits, and a named technical point of contact.

What authentication does the API use?

JWT-based token authentication. The token lifetime is a configurable parameter. When a token expires, the API returns an error; the integration should respond to this by requesting a new token and retrying — this follows the standard JWT refresh pattern. There is no separate manual renewal step, so a well-behaved integration will not stop working unattended.

What rate limits apply?

Default rate limits are sized to support routine polling and integration at the volumes typical of multi-site deployments — including high-concurrency cases involving dozens of simultaneous sessions. Customers should share their expected call patterns during onboarding so the right limits are provisioned. Specific limits and burst allowances are documented per tenant.

Is a sandbox environment with mock data available?

No, there is no separate sandbox environment with synthetic data. Integration development is done against the production API, with a Neo 1 device linked to the customer's organisation. Customers use the device to generate real test memories during the integration phase, which are then deleted before go-live if not needed. This pattern works well for most integrators because the test data exercises the same code paths as production.

Is the OpenAPI specification available as a downloadable file?

The API specification is currently provided as a Postman collection, which customers can download and import. A formal OpenAPI (JSON/YAML) export is not currently published. For most integration teams the Postman collection is sufficient to generate clients or validate request and response shapes.

Are webhooks or streaming delivery supported?

Webhooks are supported as event notifications — when a memory is created, NeoCore can notify the customer's server. However, the memory payload itself is not pushed in the webhook; the customer's server must then fetch the memory via the API. This reduces polling load but does not eliminate the API call. Customers requiring true push delivery of memory content should raise this during integration design.

What is the API versioning and deprecation policy?

A formal versioning policy is not yet published. In practice, advance notice of changes is provided two weeks to one month ahead of the change, and a period of backward compatibility is maintained for a defined window. Enterprise customers receive direct communication on any change affecting their integration.

In progress — A formal API versioning policy is being defined.

4. Reliability and Support

What uptime commitment does NeoCore offer?

Observed uptime is greater than 99%. A public status page is not currently published. For enterprise customers, a target uptime SLA is provided under the master service agreement, with measurement methodology and any service-credit structure defined in the SLA schedule.

In progress — A public status page is being planned.

What support channel is available?

Enterprise customers are supported over a combination of email, Slack, and WhatsApp, with a named technical point of contact assigned at onboarding. The scope and severity-level SLAs of support — response times, escalation paths, on-call coverage — are sized to the deployment scale and agreed during commercial discussions.

How are scheduled maintenance and incidents communicated?

Scheduled maintenance windows and active incidents are communicated through the customer's named support channel. A post-incident report is provided for material incidents affecting the customer's tenant.

What happens to memories if the API is briefly unavailable?

API availability does not affect memory synthesis as long as the device has connectivity to the NeoSapien cloud. If the customer's integration cannot reach the API for a period, memories synthesised during that window remain available on retrieval once the integration recovers — the customer can fetch them at any time using the sync endpoint. Note that this is distinct from device-side connectivity, addressed in the next section.

5. Device Filtering and Management

How are memories associated with a device or user?

Memories are associated with the account (user email) of the wearer of the device, not with a device serial number. Customers running multiple devices across a deployment should plan their data model around user identity: assign each device to a named user account, and treat the user ID as the routing key for memories.

Can the memory sync endpoint be filtered to a specific device?

Filtering is available by user ID, not device ID. To route memories to a specific physical asset, the customer maps the user wearing the device to that asset in their own data model. The mapping (user → device → asset) is maintained on the customer side; NeoCore returns the user ID with each memory and the customer's integration resolves it to the correct context.

Does each memory carry a device identifier?

No. The memory payload includes the user ID, not a device ID. See the question above for the recommended mapping approach.

Is there a device management API?

No dedicated device management API is currently exposed. User-to-device assignment is done through the admin console, and customers manage their own asset-to-user mapping externally. For enterprises running many devices, NeoSapien can support a structured onboarding process for bulk assignment.

What happens to memories captured while a device is offline?

The device requires connectivity to capture memories — there is no local memory buffer on the device itself. If the device loses connectivity, no memory is captured during the offline window. Deployments in connectivity-variable environments (vehicles, remote field operations) should plan for this and ensure a stable mobile data path during conversations the customer wants captured.

Can role-based access control be applied to who sees which users' memories?

Yes. Within a tenant, role-based access governs which administrators and analysts can view which users' memories. Visibility can be scoped to a user, a group of users, or organisation-wide — for example, restricting a regional manager to seeing only the users assigned to their region.

6. Hardware

What is the Neo 1, and what does it do?

The Neo 1 is NeoSapient's always-on, ambient listening wearable. It captures conversation in the wearer's environment and streams it to NeoCore — NeoSapient's proprietary AI platform — for transcription, diarization, and synthesis. The device is paired one-to-one with a user account; it cannot operate as a shared or unassigned asset.

How does the device connect to the cloud?

The device connects to NeoCore over mobile data, typically via a paired smartphone. No customer-side network infrastructure or IT involvement at each site is required. This makes the device suitable for deployment into the field, into vehicles, and into multi-site enterprise environments without per-site IT setup. The trade-off is that the device requires continuous connectivity to capture conversations.

Does the device buffer memories when offline?

No. There is no local memory buffer on the device — if the device is offline, conversations during that window are not captured. Customers operating in connectivity-variable environments should plan their deployment around stable mobile data coverage.

In progress — On-device buffering for offline operation is on the roadmap. Customers needing this should raise it during integration discussions.

What is the battery life?

Battery life is rated for 17 hours of continuous use on a single charge. Heavy continuous capture in noisy or high-activity environments reduces this. Specific runtime figures and the device datasheet are provided to enterprise customers during evaluation. Devices charge via USB-C.

What is the form factor?

The Neo 1 is a small pendant-style wearable designed to be worn around the neck of the user. Exact dimensions, weight, and mounting options are in the product datasheet.

What is the microphone configuration and audio capture range?

The device uses a multi-microphone array tuned for picking up the wearer's voice and nearby conversation at typical conversational distance (3 meters). Performance in highly noisy environments — moving vehicles, machinery, crowded public areas — should be validated during a pilot, since diarization accuracy reduces in those settings by approximately 4 percentage points.

Can device firmware be updated remotely?

Yes. Firmware updates are delivered over-the-air during regular cloud connectivity. Customers are notified of material firmware changes in advance through the support channel.

What is the warranty on the device?

The device carries a limited hardware warranty of twelve months from delivery against material defects, subject to the terms of the master purchase agreement. The warranty does not cover misuse, unauthorised modification, or damage outside intended operating conditions.

How many devices does a typical deployment need?

One device per active wearer. For shift-based deployments — sales territories, ambulance crews, ICU rooms, multi-shift operations — size for the maximum number of concurrent wearers, not the total headcount, since devices can be handed over between shifts (with the user account re-assignment handled in the admin console). A small spare pool is recommended for charging rotation and replacement.

What is the lead time on device orders?

Order quantities, pricing tiers, and delivery lead times are agreed under the master purchase agreement. Small orders (typically 100 devices or fewer) are dispatched on a faster cycle than larger production orders. Lead times for large enterprise rollouts are confirmed during commercial discussions.

What happens if a device is lost or stolen?

A lost or stolen device can be remotely deactivated by NeoSapien support. Deactivation prevents any further memory uploads from that device and unbinds it from the user account. Customers should report lost or stolen devices promptly to ensure timely deactivation.

Talk to us

Most enterprise pilots start with a 30-minute platform walkthrough and a scoping conversation. For platform access, pricing, or a DPA discussion:

enterprise@neosapien.xyz